



# 3D Gaussian Splatting Volumetric Pathtracing

Euan Hughes

École de Technologie Supérieure

MTR871 Presentation  
August 08, 2025

# Table of Contents

- 1 Background I: Physically-Based Volume Rendering
- 2 Background II: 3D Gaussian Splatting
- 3 Main Paper: Don't Splat Your Gaussians
- 4 Implementation
- 5 Discussion & Future Directions

# Table of Contents

- 1 Background I: Physically-Based Volume Rendering
- 2 Background II: 3D Gaussian Splatting
- 3 Main Paper: Don't Splat Your Gaussians
- 4 Implementation
- 5 Discussion & Future Directions

# What is Rendering?

**Rendering = “Light Transport Simulation”**




Image from Y. Li et al. 2019

## Physically-Based Rendering

Simulation that adheres to the laws of physics

# What is Physically-Based Volume Rendering?

**Goal:** accurately simulate light transport in complex volumetric media such as clouds, smoke, and fire.



Image from Britannica 2025



Image from Wikipedia 2025


# Why Physically-Based Volume Rendering?

**Main Application:** VFX pipelines



Image from *Up* 2009


## How does light change as it travels through a volume?






# How Physically-Based Volume Rendering?

**How does light change as it travels through a volume?**



**(!!) Must integrate over all these events**

# Volume Rendering Equation



## Transmittance

Light is continuously attenuated by extinction events. Physically described by:

$$T(\mathbf{x}_0, \mathbf{x}_t) = \exp\left(-\underbrace{\int_0^t (\mu_t(\mathbf{x}_{t'})) dt'}_{\text{'Optical Depth'}}$$

# Solving the Volume Rendering Equation

## Challenges

- Scattering happens in all directions
- Scattering happens constantly along a ray

**(!!) Especially with multiple scattering, impossible to solve analytically**

## Solution: Monte Carlo Integration!!

- Randomly sample scattering directions
- Randomly sample scattering distances

# Free-Flight Distance Sampling

## Key Insight

Transmittance formula  $T(\mathbf{x}_0, \mathbf{x}_t)$  is a '*survival function*': describes probability of a photon not interacting up to distance  $t$

## Inverse Transform Sampling

Creates a way to sample random distances consistent with this distribution

Gets a random 'target' optical depth:

$$\int_0^t (\mu_{t'}(\mathbf{x}_{t'})) dt' = \tau$$

Now, solve for distance  $t$ , which depends on extinction  $\mu_t(\mathbf{x})$

# Free-Flight Distance Sampling

**Homogeneous volumes:**  $\mu_t(\mathbf{x})$  is a fixed constant  $\mu_t$

$$\int_0^t \mu_t dt' = \mu_t t, \quad \implies \quad t = \frac{\tau}{\mu_t}$$




Image from @RaphaelRau 2022

# Free-Flight Distance Sampling

However...

Lots of volumes are *heterogeneous*: extinction varies, cannot invert CDF






Image from Kutz et al. 2017

# Regular Tracking

## Solution: Voxels!

Discretize the volume into a 3D grid of piecewise homogeneous regions.



**Then:** trace rays through the volume, track voxel boundaries, and perform homogeneous sampling over each segment.

# Regular Tracking


$$\int_0^t (\mu_t(\mathbf{x}_{t'})) dt' = \tau$$

Equivalent to:

$$\sum_{j=1}^k \mu_{t,j} \Delta_j$$

To solve for free-flight distance  $t$ :

- Sample target optical depth:  $\tau$
- Sum segment-wise depth, in order
- Stop when running sum exceeds target
- Interpolate within the final segment



**Unbiased Estimator!**



# Ray-Marching

## Regular tracking limitation

Tracking all voxel boundaries along rays are expensive, especially in very dense grids

## Alternative

March the ray in fixed spatial increments and sample whatever voxel the step occurs in, assuming homogeneous over the step.


$$\int_0^t (\mu_t(\mathbf{x}_{t'})) dt' = \tau$$

*Approximated by:*


$$\sum_{j=1}^N \mu_t(\mathbf{x}_{t'}) \Delta \quad \Rightarrow$$

**Biased Estimator!**


# Ray-Marching



**(a)** Regular Tracking Voxels



**(b)** Ray-Marching Voxels



**(c)** Ray-Marching Thin Features

# Representations

## Voxels are state-of-the-art

### Advantages

- Data can be read quickly
- Natural Level-of-Detail rendering
- Artist friendly
- Simple to integrate into pathtracers

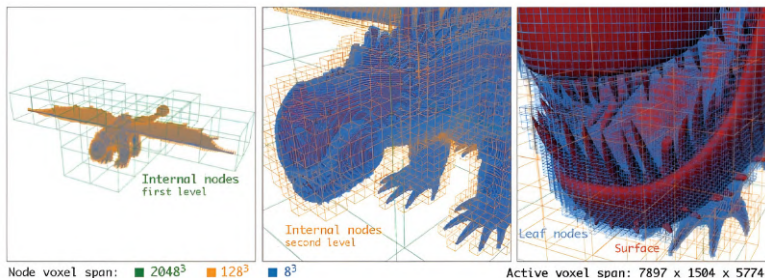



Image from Museth 2013

# Representations

## Voxel Disadvantages

- High memory requirements for dense grids
- Blocky artifacts at lower memory/Level-of-Detail



## Alternatives to voxels?

Neural representations like NeRF have impressive compression, however:

- slow lookups
- complicated to add to pathtracers
- (!! ) deviates significantly from the physical model

This motivates a new representation that:

- Adheres to the volume rendering equation
- Remains much more expressive at lower memory

# Table of Contents

- 1 Background I: Physically-Based Volume Rendering
- 2 Background II: 3D Gaussian Splatting
- 3 Main Paper: Don't Splat Your Gaussians
- 4 Implementation
- 5 Discussion & Future Directions

# 3D Reconstruction

**Goal:** infer a 3D representation of a scene from a collection of 2D images, capable of novel-view synthesis




Image from Wong 2021

# 3D Gaussian Splatting (3DGS)

## Idea:

Represent scene as collection of view-dependent Gaussians.  
Optimize parameters to match images.



Images from R. Li and Cheung 2024

**High fidelity and fast !**



# Why is 3DGS not Physically-Based?

- No light sources and scene scattering: lighting is ‘baked’ into the primitives
- “Splatting” based pipeline
  - 3D Gaussians are turned into a 2D billboard, color values are summed and alpha blended based on approximate depth ordering and a *fixed opacity*




Image from Zwicker et al. 2002

**Can we treat 3D Gaussians in a more physically-based manner? We will see: Yes !**

# The case for Physically-Based Rendering

**Primarily:** extremely expressive, superior to voxels at equivalent and even higher memory usage




Image from Zhou, Wu, and Yan 2024

**For VFX: same/higher quality renders  
with much lower memory footprint**

# The case for Physically-Based Rendering

## More desirable properties:

- Gaussians have closed form integral: the Gauss error function  
⇒ exact solution for transmittance
- inverse of Gauss error function also has a closed form  
⇒ exact solution for free-flight distances

**For VFX: we can get unbiased and physically accurate renders**

# Table of Contents

- 1 Background I: Physically-Based Volume Rendering
- 2 Background II: 3D Gaussian Splatting
- 3 Main Paper: Don't Splat Your Gaussians**
- 4 Implementation
- 5 Discussion & Future Directions

# “Gaussian-Mixture Model”




Image modified from Disney 2017

**Unlike 3DGS and prior methods:** we have light sources, and need to consider scattering events

# Gaussian-Mixture Model Formulation

**Also unlike 3DGS and prior methods:** perform exact integration over Gaussian densities

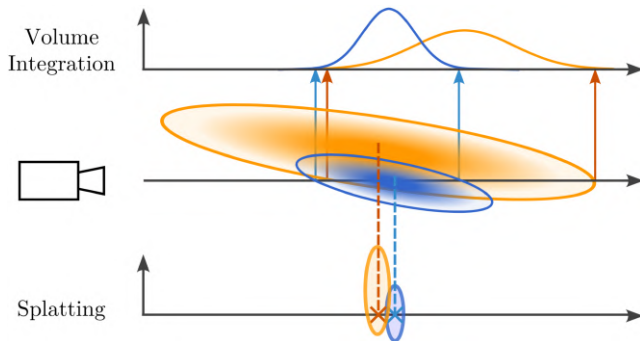


Image from Celarek et al. 2025

**We need to find exact solutions to transmittance and free-flight sampling distances over these integrals**

# Transmittance

## Extinction

$$\mu_t(\mathbf{x}) = \sum_{i=1} \sigma_i G_i(\mathbf{x})$$

**In Practice:** we have to find all entry/exit events  $[t_{i,0}, t_{i,1}]$  to find the active range of each Gaussian along a ray (we will see how in the implementation section)

## Transmittance

$$\Rightarrow T(\mathbf{x}_0, \mathbf{x}_t) = \exp\left(- \underbrace{\sum_{i=1} \sigma_i \int_{\max(0, t_{i,0})}^{\min(t, t_{i,1})} G_i(\mathbf{x}'_t) dt'}_{\text{Optical Depth}}\right)$$

$$T(\mathbf{x}_0, \mathbf{x}_t) = \exp\left(-\sum_{i=1} \sigma_i \int_{\max(0, t_{i,0})}^{\min(t, t_{i,1})} G_i(\mathbf{x}'_t) dt'\right)$$

## Solution

Since Gaussians have a closed-form integral (erf), this has an exact solution... easy :)



... harder :(

# Distance Sampling

## Ray Segmentation

- Find all Gaussian entry/exit events  $[t_{i,0}, t_{i,1}]$
- Sort all individual events (not pairs of events) in ascending order.
- This yields an ordered sequence of disjoint segments which the set of overlapping Gaussians is fixed.




Image from Condor et al. 2025

# Distance Sampling

## Recall: Regular Tracking

Accumulate optical depth over piecewise homogeneous segments until exceed target depth  $\tau$ )

**For Gaussians:** accumulate optical depth over segments with fixed numbers of Gaussians

## Problem

Once we exceed the target depth, we need to solve for the exact distance within the last segment

**In Regular Tracking:** segments are homogeneous  $\implies$  do simple linear interpolation

**For Gaussians:** ... how to solve?

## Cases:

- **One Gaussian:** closed form solution exists with the inverse error function  $\text{erf}^{-1}$
- **Multiple Gaussians:** impossible to invert analytically. Instead, must use numerical root finding method

With this formulation, we have exact solutions for transmittance and distance samples

⇒ unbiased VFX renders

**As desired:** the implementation is capable of extremely high compression rates




Image from Condor et al. 2025

# Results




Image from Condor et al. 2025

# Results



Original cloud (999x634x1224 effective resolution, 480MB VDB volume)



Ours, 16.2k primitives GMM (639KB)  
MSE:  $1.621e-4$   
Rendering time: 26.27mins



Ours, 5.6k primitives GMM (224KB)  
MSE:  $1.717e-4$   
Rendering time: 19.43mins



Ours, 2.6k primitives GMM (107KB)  
MSE:  $2.72e-4$   
Rendering time: 11.67mins

Image from Condor et al. 2025

# Table of Contents

- 1 Background I: Physically-Based Volume Rendering
- 2 Background II: 3D Gaussian Splatting
- 3 Main Paper: Don't Splat Your Gaussians
- 4 Implementation**
- 5 Discussion & Future Directions



# Building a Renderer from Scratch

Wrote full integrator from scratch

## Goals


- Proper Gaussian intersection routines
- Analytical transmittance
- Free-flight sampling

## Initial Version: Raymarching with Single Scattering

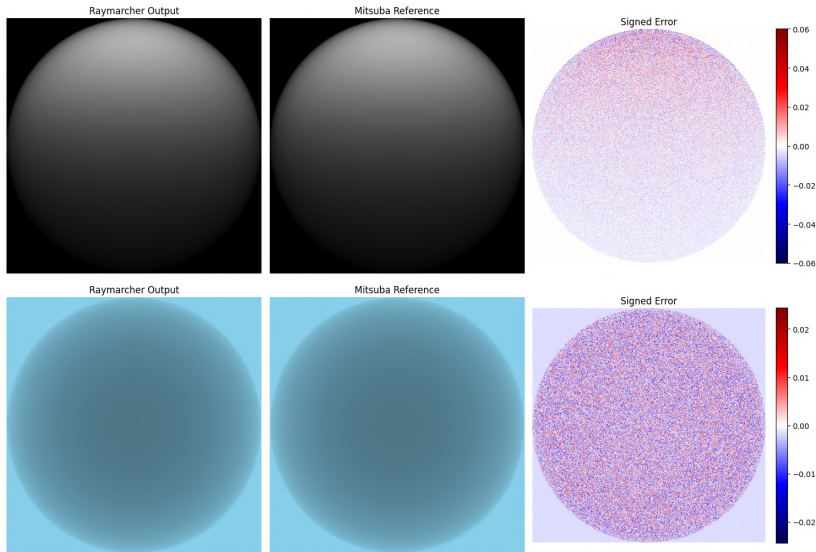
- Fixed step marching
- Compute extinction at each step  $\rightarrow$  transmittance
- In-scattered radiance from lights (also marched)
- Assumed constant extinction per step

# Validating the Raymarcher

- Tested on simple spheres
- Compared results with Mitsuba
- Matched reference images



# Validating the Raymarcher



# Representation-Agnostic Integrator

- Integrator independent of volume representation
- $\implies$  Can swap spheres  $\leftrightarrow$  Gaussians
- Useful for generating reference results for Gaussians (No Mitsuba support)

# Two Intersection Methods for Gaussians


## Method 1: Map to unit sphere

- Transform ray by whitening transform:  $\mathbf{x}' = L^{-1}(\mathbf{x} - \mu)$
- Intersect unit sphere in transformed space
- Map intersection points back to original space

## Method 2: Solve equations directly

- Treat 3D Gaussian as ellipsoid:  $(\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu) = R^2$
  - Plug ray into equation  $\rightarrow$  quadratic in  $t$
  - Solve for  $t_0, t_1$ ; clamp to ray bounds for entry/exit
- 
- Direct solver is faster
  - Both allow precomputation to speed up queries

# Validating Intersection Routine



# Extinction from Gaussians



## Recall

Extinction is sum over all Gaussians

$$\mu_t(\mathbf{x}) = \sum_i \sigma_i \cdot G_i(\mathbf{x})$$

# Reference Renders with Gaussians

- No changes to integrator
- Only volume representation changed
- Use as reference for analytical transmittance






# Adding Analytical Transmittance

- Compute exact transmittance over integral for each raymarching step
- Analytic In-Scattering
  - Find all Gaussian entry/exit points towards lights
  - Sort them
  - Compute exact transmittance to each


# Validation



## Distance Sampling Routine

- Segment primary ray at all intersections
- Sample target optical depth  $\tau$
- Accumulate  $\int \mu_t$  across segments
- When target exceeded: solve last segment (bisection)


# Validation:





# Multiple Scattering Extension

At each scattering event:

- Compute in-scattering as before
- Combine with sampling random direction and recursively walking path



# Multiple Scattering Scenes



## Volumetric Path Tracer:

- Gaussian volume support
- Multiple scattering
- Exact transmittance
- Free-flight sampling

## Next Steps:

- Acceleration structures for intersection (currently slow with many Gaussians)
- General performance optimization for large scenes

# Table of Contents

- 1 Background I: Physically-Based Volume Rendering
- 2 Background II: 3D Gaussian Splatting
- 3 Main Paper: Don't Splat Your Gaussians
- 4 Implementation
- 5 Discussion & Future Directions



# Acceleration is Crucial

## Key Finding

Rendering without acceleration structures becomes **prohibitively slow**.

Even with acceleration, most computation is spent on:

- Ray traversal
- Distance sampling
- Intersection tests

In a similar recent method *Unified Gaussian Primitives* (Zhou, Wu, and Yan 2024), this took up **85%** of the total rendering time.

# Speeding up Distance Sampling

## Distance Sampling Approximation

At the last segment, instead of solving exactly, assume homogeneity (as in raymarching).

This adds a small bias.

## Delta Tracking

Adds a majorant and performs rejection sampling.

Challenge: finding a good majorant is difficult with arbitrary overlap.

Recent work on **progressive null-collision tracking** (Misso et al. 2023) allows for high quality renders with a poor majorant.

# Speeding up Intersection

## Bounding Icosphere

Wrap each Gaussian in a bounding icosphere.

This allows fast ray-triangle intersection tests.

- Significant speedup
- But increased memory usage




Image from Condor et al. 2025

# Fighting Overlap

A major bottleneck: **high Gaussian overlap**.

## Epanechnikov Kernel

Sharper falloff  $\Rightarrow$  less overlap

- Closed-form transmittance + sampling still possible
- Often similar quality
- Much faster in practice




Image from Condor et al. 2025

# Pruning Overlap (from LoD methods)

## Idea from recent Gaussian LoD work

During optimization, prune Gaussians that overlap excessively.

- Compute average distance to  $k$  nearest neighbors
- Remove any Gaussian with excessive overlap

## Result

Speeds up rendering & reduces memory use.




Image from Seo et al. 2025

## Multi-Level Structures

Gaussians are very expressive  $\Rightarrow$  promising for LoD!

- Can adapt to memory budgets
- Scale well across detail levels

Lots of 3DGS LoD methods are agnostic to the rendering method

**But:** need to modify optimization heuristics.

## Artist Pain Point

Gaussians are much harder to author than voxels.

There must be a heuristic or method to:

- Convert voxel grids → Gaussians
- Do it **accurately** or it's useless

This is a crucial pipeline stage for real production use.

## Authors' Pipeline

Full differentiation through **stochastic multi-scattering** paths

- Much harder than 3DGS — must trace back exact random walk
- Much slower than classic 3DGS training
- But high quality and accurate

## Takeaway

Differentiability is a **must-have** for practical applications.



# Conclusion

## Big Picture

3D Gaussians show great potential in **production volumetric pathtracing**.

- Excellent visual quality
- Orders-of-magnitude lower memory usage
- Physically-based (no approximations!)

## But...

We need significant speedups in both forward and inverse rendering for widespread adoption.

**Still a young area.** Lots of promising ideas to explore in my own research.

# References I



@RaphaelRau (May 2022). *Tweet posted on May 20, 2022*. Accessed: 2025-08-07. URL: <https://x.com/OTOY/status/1527691348064669696>.



Britannica (2025). *Crepuscular ray*. Accessed 7 August 2025. Encyclopaedia Britannica. URL: <https://www.britannica.com/science/crepuscular-ray>.



Celarek, Adam et al. (Feb. 26, 2025). *Does 3D Gaussian Splatting Need Accurate Volumetric Rendering?* DOI: 10.48550/arXiv.2502.19318. arXiv: 2502.19318[cs]. URL: <http://arxiv.org/abs/2502.19318> (visited on 05/26/2025).



Condor, Jorge et al. (Feb. 28, 2025). "Don't Splat your Gaussians: Volumetric Ray-Traced Primitives for Modeling and Rendering Scattering and Emissive Media". In: *ACM Transactions on Graphics* 44.1, pp. 1–17. ISSN: 0730-0301, 1557-7368. DOI: 10.1145/3711853. URL: <https://dl.acm.org/doi/10.1145/3711853> (visited on 05/26/2025).



Disney (2017). *Walt Disney Animation Studios Cloud Data Set*. Licensed under the Creative Commons Attribution-ShareAlike 3.0 Unported License. See <http://creativecommons.org/licenses/by-sa/3.0/>. Includes volumetric cloud model photograph by Kevin Udy from the Colorado Clouds Blog (<https://coclouds.com/436/cumulus/2012-07-26/>). URL: <https://disneyanimation.com/resources/clouds/>.



Kutz, Peter et al. (Aug. 31, 2017). "Spectral and decomposition tracking for rendering heterogeneous volumes". In: *ACM Transactions on Graphics* 36.4, pp. 1–16. ISSN: 0730-0301, 1557-7368. DOI: 10.1145/3072959.3073665. URL: <https://dl.acm.org/doi/10.1145/3072959.3073665> (visited on 07/02/2025).



Li, Ruiqi and Yiu-ming Cheung (2024). "Variational multi-scale representation for estimating uncertainty in 3d gaussian splatting". In: *Advances in Neural Information Processing Systems* 37, pp. 87934–87958.



Li, Yuanhang et al. (2019). "Real-time optical 3D reconstruction based on Monte Carlo integration and recurrent CNNs denoising with the 3D light field display". In: *Optics express* 27.16, pp. 22198–22208.

# References II



Misso, Zackary et al. (2023). "Progressive null-tracking for volumetric rendering". In: *ACM SIGGRAPH 2023 conference proceedings*, pp. 1–10.



Museth, Ken (June 2013). "VDB: High-resolution sparse volumes with dynamic topology". In: *ACM Transactions on Graphics* 32.3, pp. 1–22. ISSN: 0730-0301, 1557-7368. DOI: 10.1145/2487228.2487235. URL: <https://dl.acm.org/doi/10.1145/2487228.2487235> (visited on 06/29/2025).



Seo, Yunji et al. (June 12, 2025). *FLoD: Integrating Flexible Level of Detail into 3D Gaussian Splatting for Customizable Rendering*. DOI: 10.1145/3731430. arXiv: 2408.12894 [cs]. URL: <http://arxiv.org/abs/2408.12894> (visited on 06/20/2025).



*Up* (2009). Animated feature film. United States.



Wikipedia (2025). *Deflagration*. Accessed 7 August 2025. Wikipedia, The Free Encyclopedia. URL: <https://en.wikipedia.org/w/index.php?title=Deflagration&oldid=>



Wong, Prim (Jan. 2021). *2D to 3D reconstruction*. Medium (4 min read). URL: <https://medium.com/super-ai-engineer/2d-to-3d-reconstruction-2bf4f9ede53d>.



Zhou, Yang, Songyin Wu, and Ling-Qi Yan (Sept. 22, 2024). *Unified Gaussian Primitives for Scene Representation and Rendering*. DOI: 10.48550/arXiv.2406.09733. arXiv: 2406.09733 [cs]. URL: <http://arxiv.org/abs/2406.09733> (visited on 05/26/2025).



Zwicker, Matthias et al. (2002). "EWA splatting". In: *IEEE Transactions on Visualization and Computer Graphics* 8.3, pp. 223–238.